

# Model Context Protocol (MCP) for Ansible Automation Platform

Version 1, 2026-09-26

# Welcome

## Introduction

Welcome to the learning section of this course. In this chapter, you will understand what the Model Context Protocol is, how it works, and how Ansible Automation Platform uses it to connect with AI tools.

We'll break this down into bite-sized sections that follow a simple pattern: **What is it?** → **Why is it needed?** → **How does it work?**

## What You'll Learn

By the end of this chapter, you will be able to:

- Define the Model Context Protocol and understand its role in the AI ecosystem
- Understand the architecture of MCP and how different components interact
- Explain Ansible Automation Platform's MCP server and its capabilities
- Describe the complete communication workflow from user to AI to AAP
- Identify and describe the 6 AAP MCP toolsets
- Understand the end-to-end workflow for deploying and using MCP

## Prerequisites

This course assumes that you have the following prior experience:

- Ansible Automation Platform fundamentals:
- What is AI and how to Prompt to an AI tool.

## Contributors

The PTL team acknowledges the valuable contributions of the following Red Hat associates:

- Andrew Block (Subject Matter Expert)
- Amrinder Singh (Content Architect)
- Anna Swicegood (Learning Product Manager)

# Lab Environment

## Instructions to Launch Your Lab on the Red Hat Demo Platform (RHDP)

1. **Log in** to the [RHDP portal](#).
2. **Search** for the catalog entry: **What's New in AAP 2.7: Automation Portal, MCP & EE Builder**.



### What's New in AAP 2.7: Automation Portal, MCP & EE Builder

provided by Product Technical Learning

[Order](#)  [Save as favorite](#)

**Category**  
Labs

**Provider**  
Product Technical Learning

**Product Family**  
Red Hat Automation

**Uptime**   100%

**Last update**  
[7 hours ago](#)

**Auto-Stop**  
6 Hours

**Auto-Destroy**  
48 Hours

### Description

#### Overview

The Ansible Automation Platform Portal with Ansible Automation Platform (AAP) 2.7 serves as a unified interface that shifts the platform focus from simple execution to a comprehensive "build, discover, and run" experience. This hands-on lab simplifies enterprise automation for users of all skill levels by exploring the new Automation Portal interface and the Execution Environment Builder capabilities.

Then, we have MCP. You will understand what the Model Context Protocol is, how it works, and how Ansible Automation Platform uses it to connect with AI tools.

Once the deployment is complete, you will have a guide to help you learn and explore the environment.

Duration: 1 hour

#### What You'll Learn

In this lab, you'll gain practical experience with:

- Ansible Automation Platform 2.7 with the Automation Portal interface
- Execution Environment Builder capabilities
- Defining the Model Context Protocol and understanding its role in the AI ecosystem
- Understanding the architecture of MCP and how different components interact
- Explaining Ansible Automation Platform's MCP server and its capabilities

#### Lab Environment

This lab provides a complete, hands-on environment that includes:

- Ansible Automation Platform 2.7 with Controller and Automation Hub (PAH)
- MCP and how different components interact using Open Code with Claude Code

3. **Click** on the catalog entry in the search results.
4. On the catalog page, **click** the **Order** button.
5. **Fill out** the required details in the order form.

**Activity \*** 

- ☐ Customer Facing
- ☐ Partner Facing
- ☒ Practice / Enablement

**Purpose \***

Trying out a technical solution ▼



**Salesforce ID** *(Opportunity ID, Campaign ID, CDH Party or Project ID)*

- ☐ Campaign ☐ CDH ☐ Opportunity ☐ Project 

Salesforce ID

 Search



☐ I'll provide the Salesforce ID within 48 hours. 

6. **Review the warning** at the bottom of the form and **check the box** labeled:  
“I confirm that I understand the above warnings.”
7. **Click the Order** button to place your lab order.

## Important Notes:

- This lab may take approximately **20 minutes** to become ready.
- You will receive an **email with access details** once your lab environment is ready.
- You can also **retrieve lab access** directly from the RHDP portal.

## How to Access Your Lab via the RHDP Portal:

1. On the RHDP portal, **click** on the **Services** option in the left-hand menu.
2. **Select your lab** from the listings on the right-hand side of the page to view access details.

## RHDP Portal Links

- RedHat associates: [[https://demo.redhat.com/,window=\\_blank](https://demo.redhat.com/,window=_blank)]
- RedHat partners: [[https://partner.demo.redhat.com/,window=\\_blank](https://partner.demo.redhat.com/,window=_blank)]

# Model Context Protocol (MCP) for Ansible Automation Platform

Welcome to the learning section of this course. In this chapter, you will understand what the Model Context Protocol is, how it works, and how Ansible Automation Platform uses it to connect with AI tools.

We'll break this down into bite-sized sections that follow a simple pattern: **What is it?** → **Why is it needed?** → **How does it work?**


## What You'll Learn

By the end of this chapter, you will be able to:

- Define the Model Context Protocol and understand its role in the AI ecosystem
- Understand the architecture of MCP and how different components interact
- Explain Ansible Automation Platform's MCP server and its capabilities
- Describe the complete communication workflow from user to AI to AAP
- Identify and describe the 6 AAP MCP toolsets
- Understand the end-to-end workflow for deploying and using MCP

### 1: Launching Dev Spaces from OpenShift

1. Launch the OpenShift console and login using the credentials provided on the RHDP login page.
2. Use the 9-block icon in the upper right to launch *Red Hat OpenShift Dev Spaces* directly from the OpenShift console header.




== Dev Spaces Workspace


## Launching a specific Dev Spaces Workspace

Once you have launched the OpenShift Dev Spaces console, use the following steps to access your Ansible Dev Space workspace.

1. Login via OpenShift authentication using the same credentials as the OpenShift console
2. Select **Allow selected permissions** to authorize access and launch the console
3. Scroll down and select the workspace sample titled **Ansible Dev Space**



It will take a few minutes for the Dev Spaces instance to initialize. When it has loaded, a Visual Studio Code interface will appear.




4. Click the **Trust Publishers and Install** button to trust the extension publishers that are included within the Dev Spaces workspace
5. Click the **Yes, I trust the Authors** button to trust the workspace authors (your instructors).
6. For this lab, we have provided a pre-configured configuration along with additional configurations to curate your development environment.

## Accessing a Terminal

The majority of the exercises in this lab will be performed using the Visual Studio Code terminal.

1. Open a new terminal by selecting the VS Code menu starting with the hamburger (3 horizontal lines) in the top left, then selecting **Terminal** → **New Terminal**



2. Explore the environment:

```
whoami
```

```
user
```

3. Check the Dev Spaces Workspace container OS release:

```
cat /etc/redhat-release
```

```
Red Hat Enterprise Linux release 9.7 (Plow)
```

4. This workspace provides access to **ansible** and additional developer tools:

```
ansible --version
```

```
ansible [core 2.20.5]
  config file = None
  configured module search path = ['/projects/.ansible/plugins/modules',
'/usr/share/ansible/plugins/modules']
  ansible python module location = /usr/local/lib/python3.12/site-packages/ansible
```

```
ansible collection location =  
/projects/.ansible/collections:/usr/share/ansible/collections  
executable location = /usr/local/bin/ansible  
python version = 3.12.12 (main, Apr 16 2026, 00:00:00) [GCC 11.5.0 20240719 (Red  
Hat 11.5.0-11)] (/usr/bin/python3.12)  
jinja version = 3.1.6  
pyyaml version = 6.0.1 (with libyaml v0.2.5)
```

**Ready?** Start with Section 1 below!

## What is MCP?

### What is the Model Context Protocol?

The **Model Context Protocol (MCP)** is an open-source standard that provides a universal way for AI models to interact with external data and tools.

Think of it as "**USB-C for the AI ecosystem**" — just like USB-C works with many different devices without custom adapters, MCP works with many different AI models and tools without requiring custom integrations.

[MCP as USB-C for AI] | *mcp-concept.png*

Instead of building a unique integration every time you want an AI model to interact with a new system, MCP provides a standardized method. This means:

- AI models can connect to different data sources without model-specific code
- You can swap AI providers without rewriting integrations
- Security and permissions are handled consistently

### Why is MCP Needed?

Imagine you want to use ChatGPT to query your servers, while a colleague uses Claude to do the same thing. Without a standard protocol:

- ChatGPT integration team writes custom code to connect ChatGPT to your servers
- Claude integration team writes completely different code to do the same thing
- You have duplicated code, duplicated bugs, and duplicated maintenance

**With MCP**, both ChatGPT and Claude connect to the same MCP server. No duplication. No redundancy.

### The Real Problem MCP Solves

AI models like Claude or ChatGPT are trained on data **frozen at a point in time**. They don't know:

- What jobs ran in your Ansible Automation Platform **yesterday**

- Which servers are currently online in your inventory
- What security events happened **today**
- Real-time metrics or current status

MCP solves this by giving AI models **real-time access** to your live systems. They can ask questions and get current answers, not guesses based on training data.

## How Does It Work? (Simple Version)

MCP works in three simple steps:

### *MCP Flow*

1. **You ask** → You send a natural language request to your AI tool
2. **AI translates** → The AI understands your intent and calls the appropriate MCP tool
3. **MCP delivers** → The MCP server fetches real data and returns it to the AI

### *Example*

```
You: "List my recent Ansible jobs"
  []
AI: "I need to call the job_list tool with recent=true"
  []
MCP Server: Queries AAP and returns the list of jobs
  []
AI: "Here are your last 5 jobs: ..."
```

## Next Steps

You understand **what** MCP is. Now let's understand **how** it's built — the architecture and components.

# MCP Architecture

## What are the Building Blocks?

MCP is built from **four main components**. Let's break them down:

### Component 1: The AI Client (The Requester)

This is where **you** interact with the AI. It could be:

- **Claude Desktop** — A desktop application by Anthropic
- **Cursor IDE** — A code editor with AI built in
- **LibreChat** — A web-based chat interface
- **ChatGPT** — OpenAI's chatbot

- **Any MCP-compatible tool** you're using

**Role:** You enter natural language requests here. "Show me failed jobs from yesterday." The client sends your request to the AI model.

## Component 2: The AI Model (The Translator)

This is the **large language model (LLM)** — the brain of the operation.

Examples: \* Claude (Anthropic) \* GPT-4 (OpenAI) \* Qwen (Alibaba) \* Any LLM that supports MCP

**Role:** It interprets your natural language request and figures out which MCP **tools** to call and with what parameters.

### *Example Interpretation*

Your Input: "How many Windows servers do I have?"

AI Model Thinks: "The user wants server counts. I should:

1. Call `inventory_list` to get hosts
2. Filter for hosts with 'os\_windows' group
3. Count them"

## Component 3: MCP Server (The Gatekeeper)

The **MCP Server** is the bridge between the AI and your infrastructure.

**Role:** \* Receives structured requests from the AI model \* Validates the request using your API token \* Executes the request against your system (AAP, in this case) \* Standardizes the response and sends it back to the AI

The MCP server **never trusts the AI directly**. It requires authentication and applies your system's security policies.

## Component 4: Your System (The Executor)

In this course, it's **Ansible Automation Platform (AAP)**, but it could be:

- Your cloud platform (AWS, Azure, GCP)
- Your internal database
- Your monitoring system
- Any system with an API


## Component 5: The Response Loop (The Feedback)

Response loop: The automation result is returned to the MCP server, standardized into a format the AI agent can process, and displayed to the user via the AI client.

**Role:** Executes the actual action or returns the requested data.

## How They Connect

Here's the complete architecture:



## What About Security?

The MCP Server is the **security checkpoint**:

1. **Authentication**: Uses your AAP API token to verify requests
2. **Authorization**: Only allows actions your user account has permission to perform
3. **Standardization**: Ensures all requests follow the same format
4. **Validation**: Checks that parameters are valid before sending to AAP

### Security Example

AI asks: "Delete all hosts"

MCP Server checks:

- ✓ Is the API token valid? YES
- ✓ Does the user have delete permission? NO
- ✗ DENIED - Permission rejected

Response to AI: "You don't have permission to delete hosts"

## Next Steps

You understand the components. Now let's see how **Ansible Automation Platform** uses MCP.

## AAP Integration & Toolsets

### What is the Ansible MCP Server?

The **Ansible MCP Server** is Ansible Automation Platform's implementation of the MCP standard.

It enables connecting AAP with any preferred external AI tool (such as Claude, Cursor, or ChatGPT). The AI tools can access information about your AAP environment and perform tasks using natural language.

Instead of logging into the AAP web UI manually, you can ask an AI: "Launch the 'Deploy Web Servers' job template" and it does it for you.

### Why AAP Needs an MCP Server

AAP manages:

- **Thousands of servers and devices**
- **Complex automation workflows and job executions**
- **Sensitive credentials and access control**
- **User teams and role-based permissions**
- **Audit logs and compliance data**

Before MCP, the only way to query this was:

1. Log into the web UI manually
2. Click through menus
3. Run API commands with `curl` or write custom scripts

**MCP enables:** Ask the AI in natural language, get instant answers.

## The 6 AAP MCP Toolsets

The Ansible MCP Server exposes **6 separate toolsets**, each focused on a different area of AAP:

### 1. Job Management

**What it does:** Execute and monitor automation jobs.

**Tools you can call:**

- List job templates
- Start a job run
- Check job status
- View job output/logs
- List recent jobs

**Example queries:**

```
"Show me failed jobs from the last 24 hours"  
"Launch the 'Configure Servers' job template"  
"What jobs are currently running?"
```

### 2. Inventory Management

**What it does:** Manage hosts, groups, and inventories.

**Tools you can call:**

- List inventories
- Get hosts in a group
- List inventory groups
- Get host details and variables
- Filter hosts by properties

**Example queries:**

```
"Show me all Windows hosts"
```

```
"Which hosts are in the 'production' group?"  
"List hosts that haven't checked in for 7 days"
```

### 3. System Monitoring

**What it does:** Monitor platform health and performance metrics.

**Tools you can call:**

- Check AAP health status
- Get execution node metrics
- View job execution statistics
- Monitor capacity and usage
- Get system performance data

**Example queries:**

```
"What's the health status of my AAP instance?"  
"Show me job execution metrics for the past week"  
"What's the capacity utilization?"
```

### 4. User Management

**What it does:** Manage users, teams, and organizations.

**Tools you can call:**

- List users and teams
- Get team memberships
- View organization details
- Check user roles and permissions
- Manage RBAC assignments

**Example queries:**

```
"List all users in my organization"  
"Who are the members of the ops team?"  
"What permissions does user 'alice' have?"
```

### 5. Security & Compliance

**What it does:** Access security policies, credentials, and audit logs.

**Tools you can call:**

- View audit logs
- List credentials
- Check security settings
- Review compliance data
- Monitor access and changes

#### Example queries:

```
"Show me audit logs from the past 24 hours"
"What credentials are stored in AAP?"
"List all security-related activity"
```

## 6. Platform Configuration

**What it does:** Configure platform settings and resources.

#### Tools you can call:

- List projects and their status
- View execution environments
- Get platform settings
- Manage integrations
- Configure system properties

#### Example queries:

```
"List all projects and their SCM status"
"What execution environments are available?"
"Show me platform configuration settings"
```

## One Server, Many Endpoints

All 6 toolsets are served by **one MCP server** with **6 separate endpoints**:

| Toolset                | Endpoint Path               |
|------------------------|-----------------------------|
| Job Management         | /job_management/mcp         |
| Inventory Management   | /inventory_management/mcp   |
| System Monitoring      | /system_monitoring/mcp      |
| User Management        | /user_management/mcp        |
| Security & Compliance  | /security_compliance/mcp    |
| Platform Configuration | /platform_configuration/mcp |

When you configure an AI tool, you connect it to **all 6 endpoints** so the AI has access to everything.

## Role-Based Access Control (RBAC)

The MCP server respects your AAP permissions. If you create an API token for a user with "read-only" permissions, the AI can only read data—it cannot launch jobs or make changes.

### *RBAC Example*

Token created with SCOPE: Read

Allowed:

- ✓ List jobs
- ✓ View job output
- ✓ List inventories
- ✓ View hosts

Denied:

- ✗ Launch a job
- ✗ Create users
- ✗ Delete credentials
- ✗ Modify settings

This means:

- **Admins** can use the AI to do anything they could do manually
- **Read-only users** can only query information
- **Team members** respect team-level permissions
- **Audit logs** track all AI-driven actions

## Next Steps

You understand the toolsets. Now let's **Know More About MCP Tool Discovery and Access Control**.

# Know More About MCP Tool Discovery and Access Control

## Optimize Token Usage with the Discover Endpoint

When integrating an AI agent with an MCP server, loading every tool for each request can increase token consumption. The discover endpoint helps optimize this by providing a lightweight summary of available toolsets.

The AI agent identifies the relevant toolset and loads only the tools required for the task.

```
<MCP server URL>/mcp/discover
```

Once loaded, tools remain available for the rest of the session. Additional toolsets are loaded as needed.

## Access All Tools

If token optimization is not a priority, use the all-tools endpoint to make every tool available.

```
<MCP server URL>/mcp
```

## Restrict Access to Specific Toolsets

To control which tools the AI agent can access, configure a separate endpoint for each toolset.

```
<MCP server URL>/mcp/<toolset>
```

Available toolsets: `job_management`, `inventory_management`, `system_monitoring`, `user_management`, `security_compliance`, `platform_configuration`, `content_discovery`.

## Understand Server-Level and User-Level Permissions

The MCP server uses two layers of access control:

- **Server-level permissions:** Configured by the administrator during deployment.
  - Read-only: Allows the AI agent to retrieve information without making changes.
  - Read-write: Allows changes, subject to the user's API token permissions.
- **User-level permissions:** Inherited from the user's Ansible Automation Platform API token. The AI agent can perform only actions authorized for that user.

For example, if a user can access the `network` inventory but not the `database` inventory, the AI agent cannot access the database inventory using that user's token.

## Security Consideration

Read-write access allows the AI agent to make changes in Ansible Automation Platform, including launching automation jobs when authorized.

Because AI agents may misinterpret requests, start with read-only access and introduce write permissions only when appropriate safeguards and RBAC controls are in place.

## End-to-End Workflow

## Getting MCP Running: The Three Steps

Now that you understand **what** MCP is and **how** it works, here's what you actually **do** to make it work and connect your external AI tool to Ansible Automation Platform (AAP) through the Ansible MCP Server. The process is broken down into three steps:

### Requirements to Deploy the MCP Server

To make use of the Ansible MCP Server, the following steps are required:

1. Deploy the Ansible MCP Server within Ansible Automation Platform
2. Generate an API token for use by the external AI tool
3. Configure the external AI tool to connect to the Ansible MCP Server using the generated API token

### Step 1: Deploy the MCP Server

#### Where Does It Run?

The Ansible MCP Server runs as a containerized/pod service on your infrastructure. You can deploy it on:

- **OpenShift Container Platform (OCP)** — Most common for enterprise
- **RHEL with containers** — Simpler, on a single machine
- **Kubernetes** — Any Kubernetes cluster

Deploy the Ansible MCP Server to your OpenShift environment by enabling the MCP capability within the `AnsibleAutomationPlatform` custom resource.

1. Launch the OpenShift Web Console from the RHDP login page and use the credentials provided to login to the OpenShift console.
2. Navigate to **Ecosystem** → **Installed Operators**
3. From the *Project* dropdown, ensure `aap` is selected
4. The `Ansible Automation Platform` operator is listed and installed within this project.

Red Hat OpenShift

Project: aap


### Installed Operators

Installed Operators are represented by ClusterServiceVersions within this Namespace. For more information, see the [Understanding Operators documentation](#). Or create an Operator and ClusterServiceVersion using the [Operator SDK](#).

Name Search by name...

| Name                                                                                                                                                                             | Managed Namespaces                                                                                                                                                               | Status                                                                                                    | Last updated                                                                                             | Provided APIs                                                                                                                                                                            |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|  <b>Ansible Automation Platform</b><br>2.6.0-0.1767630689 provided by Red Hat                   |  <b>NS aap</b>                                                                                  |  Succeeded<br>Up to date |  Jan 8, 2026, 2:32 PM | <a href="#">Automation Controller Backup</a><br><a href="#">Automation Controller Restore</a><br><a href="#">Automation Controller Automation Hub</a><br><a href="#">View 20 more...</a> |
|  <b>Streams for Apache Kafka</b><br>3.1.0-6 provided by Red Hat                                 |  <b>NS aap</b><br>The operator is running in openshift-operators but is managing this namespace |  Succeeded<br>Up to date |  Jan 8, 2026, 2:12 AM | <a href="#">Kafka</a><br><a href="#">Kafka Node Pool</a><br><a href="#">Kafka Connect</a><br><a href="#">Kafka Bridge</a><br><a href="#">View 5 more...</a>                              |
|  <b>Red Hat OpenShift Dev Spaces</b><br>3.25.0 provided by Red Hat                              |  <b>NS aap</b><br>The operator is running in devspaces but is managing this namespace           |  Succeeded<br>Up to date |  Jan 8, 2026, 2:10 AM | <a href="#">Red Hat OpenShift Dev Spaces instance Specification</a>                                                                                                                      |
|  <b>DevWorkspace Operator</b><br>0.38.0 provided by Devfile                                     |  <b>NS aap</b><br>The operator is running in devspaces but is managing this namespace           |  Succeeded<br>Up to date |  Jan 8, 2026, 2:09 AM | <a href="#">DevWorkspace</a><br><a href="#">DevWorkspaceTemplate</a><br><a href="#">DevWorkspaceOperatorConfig</a>                                                                       |
|  <b>Gitea Operator</b><br>2.1.0 provided by Red Hat Portfolio Technical Marketing and Platforms |  <b>NS aap</b><br>The operator is running in gitea-operator but is managing this namespace      |  Succeeded<br>Up to date |  Jan 8, 2026, 1:49 AM | <a href="#">Gitea</a>                                                                                                                                                                    |
|  <b>Red Hat OpenShift GitOps</b>                                                                |  <b>NS aap</b><br>The operator is running in openshift-operators                                |  Succeeded<br>Up to date |  Jan 8, 2026, 1:48 AM | <a href="#">Argo CD NamespaceManagement</a>                                                                                                                                              |

- Click on the **Ansible Automation Platform** operator which will display all of the resources that are managed by the operator.
- In the top tab, click on **Ansible Automation Platform** and select **aap**.



7. Click on the **YAML** tab to edit the resource in YAML format.

8. Add the following to the **spec** section by updating the configuration to enable the MCP server:

```
...
spec:
  mcp:
    disabled: false
    allow_write_operations: false
```



Setting **allow\_write\_operations** to **false** will restrict the MCP server to read-only operations. This is the preferred option, especially in Enterprise environments, when beginning to incorporate AAP with external AI tools. The AI tool may perform unwanted actions if write operations are enabled without fully understanding the implications. For the purpose of this lab, access to write operations will be restricted.

9. Click the **Save** button to apply the changes.

Once the configuration has been applied, the Ansible Automation Platform operator will begin to reconcile the changes. This process can take several minutes to complete.

The OpenShift Operator will then:

1. Create a new Pod for the MCP server

2. Expose it with a network route
3. Handle certificates and networking

## How Long Does It Take?


Usually **5-15 minutes** for the operator to reconcile and the MCP server to be ready.

## Step 2: Generate an API Token


### What's an API Token?

An **API token** is like a password, Create a Personal Access Token (PAT) by performing the following:

1. Launch the Ansible Automation Platform web interface and login using the credentials from the RHDP login page.
2. From the navigation menu on the left, expand **Access Management** and click on **API Token**.



3. On the API Token page, click the **Create API Token** button
4. On the Create API Token page, enter the following details:
  - **Description:** mcp
  - **OAuth Application:** Leave this field blank which will result in a Personal Access Token (PAT) being created.
  - **Scope:** Select **Read** to restrict the token to read-only operations.



5. Click **Create token** to generate the API token which will be displayed in the resulting dialog



Make sure to store the token securely, as it will not be displayed again.

### 3: Configuring External AI Tools

Once the Ansible MCP Server is deployed and the API token is generated, the next step is to configure an external AI tool to connect to the MCP server.

Several AI tools can be used for this integration.

- link:[[OpenCode](#)](OpenCode) — A terminal-based AI agent that can run from an OpenShift Dev Spaces workspace.
- link:[[LibreChat](#)](LibreChat) — A web-based AI chat tool that can be accessed through a browser.

Other AI tools, such as ChatGPT or Claude, may also support MCP integrations depending on their current capabilities and configuration options.

In the following sections, you will learn how to configure **OpenCode to connect to the Ansible MCP Server** using the generated API token where we will be using the claude-sonnet-4-6 model from the LiteLLM inference service to process natural language requests and interact with Ansible Automation Platform through the Ansible MCP Server.

#### 3.1: Using OpenCode with the Ansible MCP Server

[OpenCode](#) is an Open Source, terminal based AI agent. It is already installed in your Dev Spaces workspace and requires that a few steps be completed prior to its use.

### 3.1.1: Customizing a Dev Spaces Workspace

OpenShift Dev Spaces provides several ways to customize a workspace. In this section, you will create a ConfigMap and Secret in the admin-devspaces project to store and inject environment variables into the workspace. These values will be referenced by the OpenCode configuration file used to connect to the Ansible MCP Server.

To make the resources available inside the Dev Spaces workspace:

- Add the **controller.devfile.io/mount-to-devworkspace: "true"** label to the ConfigMap and Secret.



The value of the controller.devfile.io/mount-as annotation in the snippet above is set to env, which means the contents of the ConfigMap will be mounted as environment variables in the Dev Spaces workspace.

The Dev Spaces controller automatically detects and mounts resources with the required label.


First, create a **ConfigMap** resource to contain the non-sensitive configuration values by performing the following in the OpenShift web console. Instead of using the OpenShift web console to interactively create the ConfigMap, you will import the raw YAML definition using the **Import YAML** feature.

1. Launch the OpenShift Web Console from the RHDP login page and use the credentials provided to login to the OpenShift console.
2. On the top navigation bar, click on the + icon and select **Import YAML**

The screenshot displays the OpenShift web console interface. The top navigation bar includes a menu icon, the Red Hat OpenShift logo, and a user profile dropdown labeled 'admin'. A yellow box highlights a '+' icon in the navigation bar, with a dropdown menu showing options: 'Import YAML', 'Import from Git', and 'Container images'. A tooltip above the menu reads 'Create resources from their YAML or JSON definitions'. The main content area shows the 'Overview' page for a cluster, with sections for 'Getting started resources', 'Details', 'Status', and 'Activity'. The 'Status' section shows a green checkmark for the 'Cluster' and a yellow arrow pointing to the 'Control Plane'.

3. Enter the following into the *Import YAML* dialog:

```
kind: ConfigMap
apiVersion: v1
metadata:
  name: mcp-devspaces-config
  namespace: admin-devspaces
  labels:
    controller.devfile.io/mount-to-devworkspace: 'true'
    controller.devfile.io/watch-configmap: 'true'
  annotations:
    controller.devfile.io/mount-as: env
data:
  AAP_MCP_URL: https://<aap-mcp-instance-URL>/mcp
  LITELLM_MODEL_NAME: openai/granite-3-2-8b-instruct, claude-sonnet-4-6, gemini-2.5-pro
  LITELLM_URL: <LLM base URL from RHDP login page>
  OPENCODE_CONFIG: /projects/ansible-devspaces/opencode.json
```



Notice the value of the `controller.devfile.io/mount-as` annotation in the snippet above is set to `env`, which means the contents of the `ConfigMap` will be mounted as environment variables in the Dev Spaces workspace and for getting the `AAP_MCP_URL` value, you can go to OCP console → Networking → Routes and look for the route that is created for the MCP server. The URL will


be in the format of `https://aap-mcp-<instance_name>.<cluster_domain>/mcp` where `<instance_name>` is the name of your Ansible Automation Platform instance and `<cluster_domain>` is the domain of your OpenShift cluster.

4. Click the **Create** button to create the ConfigMap resource.

Now, create a **Secret** resource to contain the sensitive configuration values by repeating the same steps as above by using the *Import YAML* feature.

1. On the top navigation bar, click on the + icon and select **Import YAML**
2. Enter the following into the *Import YAML* dialog. Make sure to replace `PLACEHOLDER_REPLACE_ME` with the AAP API Token you generated in the prior section:


```
kind: Secret
apiVersion: v1
metadata:
  name: mcp-devspaces-config
  namespace: admin-devspaces
  labels:
    controller.devfile.io/mount-to-devworkspace: 'true'
    controller.devfile.io/watch-secret: 'true'
  annotations:
    controller.devfile.io/mount-as: env
stringData:
  AAP_MCP_API_KEY: <Replace this with the API Token generated from the Ansible
Automation Platform>
  LITELLM_API_KEY: <API Key available from the LLM Information for the MCP Course
RHDP login page>
type: Opaque
```




3. Click the **Create** button to create the Secret resource.

With the environment variables defined in the **ConfigMap** and **Secret** now defined within OpenShift, the next step is to restart your Dev Spaces workspace so that it can mount the values contained within these resources. To restart your Dev Spaces workspace, perform the following steps in the OpenShift web console:

1. Use the 9-block icon in the upper right to launch Red Hat OpenShift Dev Spaces directly from the OpenShift console header.



2. Within the Dev Spaces console, on the *Workspaces* page, click the kebab menu (three vertical dots) next to your `ansible-dev-spaces-example-xxxx` workspace and select **Restart Workspace**.



3. Click the **Open** link next to your workspace to relaunch the workspace

When the workspace has started, you can verify that the environment variables are correctly mounted by executing the following command in the terminal which will check for the presence of the environment variables as defined in the **ConfigMap** and **Secret** resources:

```
env | grep -E "(AAP_MCP|LITELLM)"
```

The output should display values that correspond to the environment variables defined in the **ConfigMap** and **Secret** resources previously.

### 3.1.2: Configuring OpenCode

With the required environment variables mounted in your Dev Spaces workspace, the next step is to configure OpenCode to interact with Ansible Automation Platform through the Ansible MCP Server.

OpenCode is an open-source AI development platform that provides a unified interface for working with AI models and services.

Earlier, you configured the `OPENCODE_CONFIG` environment variable in the **ConfigMap**. This variable specifies the location of the OpenCode configuration file.

Create a new file named **opencode.json** at the within the `/projects/ansible-devspaces` directory containing the following:

*/projects/ansible-devspaces/opencode.json*

```
{
  "$schema": "https://opencode.ai/config.json",
  "autoupdate": false,
  "enabled_providers": [
    "LiteLLM"
  ],
  "mcp": {
    "aap": {
      "type": "remote",
      "enabled": true,
      "oauth": false,
      "url": "{env:AAP_MCP_URL}",
      "headers": {
        "Authorization": "Bearer {env:AAP_MCP_API_KEY}"
      }
    }
  },
  "provider": {
    "LiteLLM": {
      "npm": "@ai-sdk/openai-compatible",
      "name": "LiteLLM",
      "options": {
        "baseUrl": "{env:LITELLM_URL}",
        "apiKey": "{env:LITELLM_API_KEY}"
      }
    }
  }
}
```

```

    },
    "models": {
      "claude-sonnet-4-6": {
        "name": "{env:LITELLM_MODEL_NAME}",
        "limit": {
          "context": 35000,
          "output": 1000
        }
      }
    }
  }
}

```

In this configuration file, we are specifying details related to our inference service along with the AAP MCP Server. Take note of the variables beginning with `{env:…}` which will replace the values with the content injected from the *Secret* or *ConfigMap* defined previously.

### 3.1.3: Using OpenCode to Interact with the Ansible MCP Server

With all of the configuration steps complete, let's utilize OpenCode to interact with the Ansible MCP Server.

To launch OpenCode, execute the following command:

```
opencode
```

Like many other AI coding tools, OpenCode is a Text User Interface (TUI) based tool which allows you to interact directly with the tool. By launching OpenCode, you will be placed at a prompt for which you can enter commands or execute queries:



As you can see, OpenCode is using the LiteLLM inference service along with the `{litellm_available_models_list}` Model that was defined in the *ConfigMap*.

Use the `/status` command to view the status of the tool including any integration with external components, such as the AAP MCP server.

```
/status
```

If your MCP server is running and the API token that was stored in the *Secret* is being loaded properly, the status being displayed should look similar to the following:



The **aap** MCP server displays *Connected* indicating that it is able to integrate with Ansible Automation Platform as needed. If the status being displayed indicates an error, investigate any message/error code that may be present. In addition, you can view the logs from the AAP MCP Server pod that is running in the **aap** Project. The pod will have a name of `<instance_name>-mcp-xxxx` which may provide additional insights related to the failed integration.

Hit the **ESC** key to exit the status dialog to return to the main prompt.


Now, let's start a chat session using our natural language that retrieves details that originate from Ansible Automation Platform that is facilitated through the integration with the AAP MCP Server.

Let's ask for information related to the *Projects* that have been defined:

Display the names of aap projects

In a few moments, you will see OpenCode making queries to the *Tools* exposed by the AAP MCP Server and then providing a natural language response based on the information retrieved.

A successful invocation of the **project\_list** MCP tool will display similar to the following:



Either the list of AAP project that are returned as a result of the **projects\_list** tool invocation will be displayed or you can continue the conversation by asking additional questions to display more information related to the desired query.



Due to the non-deterministic nature of LLM's, the exact response provided by OpenCode may vary between invocations and may not match exactly what is

shown in the screenshot or steps above.

Feel free to continue the chat session or start a new session to continuing to converse with the LLM, obtaining realtime insights from Ansible Automation Platform though the use of the Ansible MCP Server.

## Key Takeaways

- **MCP provides a standard way** to connect AI models with external systems.
- **Ansible MCP Server connects AI with AAP**, exposing AAP capabilities through MCP toolsets.
- **The MCP Server acts as the gatekeeper**, handling requests between the AI client and AAP.
- **Authentication and RBAC are enforced**, ensuring users can only access and perform actions permitted by their AAP permissions.
- **AAP data is standardized**, providing the AI with consistent and easier-to-process responses.
- **Natural-language interaction simplifies automation**, reducing the need for manual UI navigation or direct API calls.

# Error Detection & AI-Powered Problem Solving with MCP

Welcome to Chapter 2! In Chapter 1, you learned **what** MCP is and **why** it matters. Now, you'll learn how MCP and Claude AI work together to solve Ansible playbook errors **automatically**.

We'll break this down into bite-sized sections that follow a simple pattern: **What is the problem?** → **How does MCP solve it?** → **What does this look like in practice?**

## What You'll Learn

By the end of this chapter, you will be able to:

- Understand how Ansible playbooks fail and the typical troubleshooting process
- Explain how MCP enables automatic error detection and resolution
- Describe the complete workflow from error occurrence to fix application
- Understand how Claude AI analyzes errors and generates solutions
- Implement the operator approval workflow for safety and control
- Troubleshoot issues in the error resolution system

## The Big Picture: Self-Healing Playbooks

Imagine a world where:

- Your Ansible playbook fails
- The system **automatically captures** the error details
- Claude AI **analyzes** the error in seconds
- A Human **reviews** Claude's suggested fix and implement it.
- Your playbook **retries** and succeeds

This is what we'll build in this chapter.

## The Problem We're Solving

### Current Situation (Without MCP)

When a playbook fails, teams typically:

1. Get alerted to the error.
2. Search for a solution.
3. Diagnose the problem.
4. Develop and test a fix.

5. Apply the fix to production.
6. Verify the result.

This process is **manual, reactive, and time-consuming**, often taking 1–2 hours or longer.

## New Situation (With MCP + Claude)

With MCP and Claude:

1. The error is automatically captured with context.
2. **Claude analyzes the error and suggests a solution.**
3. **A human reviews and implements the suggested solution.**
4. The playbook is rerun and the result is verified.

This approach makes troubleshooting **faster and more efficient while keeping the human in control of the fix.**

## Prerequisites for This Chapter

To get the most from this chapter, you should:

- □ Have completed Chapter 1 (understand MCP basics)
- □ Understand basic Ansible playbook concepts
- □ Know what an Ansible error looks like
- □ Have access to AAP or test Ansible environment
- □ Have a Claude API key (from anthropic.com)

## Mcp Server for playbook error detection and resolution.

This brief will demonstrate how the MCP server can be leveraged to identify errors in Ansible playbooks, use Claude to analyze the issues and recommend appropriate resolutions, and then automate the remediation process. To ensure the changes are applied in a controlled and governed manner, the Automation Orchestrator approval workflow can be integrated to provide a human-in-the-loop validation step before the fix is automatically implemented.

### Lab Scenario: Hello World Sample with Error

We'll use a simple but intentional error to demonstrate the system:

```
---
- name: Hello World Sample
  hosts: all
  tasks:
```

```
- name: Hello Message
  debug:
    msg: "Hello World"

- name: APP version
  debug:
    msg: "{{ aap_version }}"
```

we created a file `hello-world-playbook.yml` with the above content. The second task references a variable `aap_version` that is not defined anywhere, which will trigger an error when the playbook is executed. This will allow us to observe how MCP captures the error and a prompt to Claude analyzes it.

## Part 1: Configure GitHub in Ansible Automation Platform

### Step 1.1: Access AAP Controller

1. Navigate to your AAP Controller URL from the RHDP login page.
2. Log in with your credentials.

### Step 1.2: Create Project in AAP

1. Go to **Projects** (Resources → Projects)
2. Click **Create project**
3. Fill in the project details:

| Field               | Value                                                                                                                                                     |
|---------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------|
| Name                | MCP-Playbooks-Lab                                                                                                                                         |
| Description         | "Ansible Playbooks with MCP Error Detection"                                                                                                              |
| Source Control Type | Git                                                                                                                                                       |
| Source Control URL  | <a href="https://github.com/opentraining-labs-customizations/aap-playbooks.git">https://github.com/opentraining-labs-customizations/aap-playbooks.git</a> |

4. Click **Save**
5. Wait for the project to sync (you'll see a status message)

```
Status: Successful
Commit: abc1234567...
Branch: main
```

📌 **GitHub is now connected to AAP!**

## Part 2: Create Job Template with MCP Integration

In AAP Controller:

1. Go to **Templates** (Resources → Templates)
2. Click **Create template** → **Job template**
3. Fill in the basic details:

| Field       | Value                                           |
|-------------|-------------------------------------------------|
| Name        | MCP-Error-Detection-Lab                         |
| Description | "Hello World playbook with MCP error detection" |
| Job Type    | Run                                             |
| Inventory   | Demo-inventory                                  |
| Project     | MCP-Playbooks-Lab                               |

| Field    | Value                    |
|----------|--------------------------|
| Playbook | hello-world-playbook.yml |


4. Click **Save**

☐ **Job Template is ready!**

## Part 3: Run the Playbook

In AAP Controller:

1. Go to **Templates** (Resources → Templates)
2. Find **MCP-Error-Detection-Lab**
3. Click the **Launch** button (play icon)



### Step 3.3: Monitor Job Execution

The job output will show:

```
PLAY [Hello World Sample] ****

TASK [Hello Message] ****
ok: [localhost] => {
    "msg": "Hello World"
}

TASK [APP version] ****
fatal: [localhost]: FAILED! => {
    "msg": "undefined variable: aap_version"
}

PLAY RECAP ****
localhost: ok=1 changed=0 unreachable=0 failed=1


FAILED - RETRYING: APP version (Attempt 1 of 3)
```

+ image::1.4.png[]

❑ **Error detected! Now lets check with the MCP Server**


### Observe MCP Error Detection & Resolution

1. Go to Ansible devspace and open the vscode terminal and Run the command `opencode` to start the MCP server. You should see the following output:
2. Now give a Prompt to Claude to analyze the error. The system will automatically send the error context to Claude for analysis. Within 30 seconds, you will receive a proposed solution.



```
why my MCP-Error-Detection-Lab job template is failing help me resolve this error ?
```

3. The output can be different from the following example, but it should look similar and the resolution would be to add the `aap_version`.



```
* aap_job_templates_list [search=MCP-Error-Detection-Lab]

Build · openai/granite-3-2-8b-instruct,claude-sonnet-4-6,gemini-2.5-pro
```

```

Objective
- User wants to diagnose and resolve the failure in the MCP-Error-Detection-Lab job template in AAP.

Important Details
- Job template ID: 8, name: MCP-Error-Detection-Lab
- Last job run: job ID 2, status: failed
- Project ID: 7, Inventory: Demo Inventory (ID: 1)
- Inventory has active failures: has_active_failures: true, hosts_with_active_failures: 1
- Playbook: hello-world-playbook.yml
- Execution environment: Default execution environment (ID: 2)

Work State

Completed
- Listed AAP projects (1 project: Demo Project, ID: 5)
- Retrieved job template details for MCP-Error-Detection-Lab (ID: 8)
- Retrieved job run details for job ID 2
- Retrieved job stdout - root cause identified:
  - Task APP version fails with: fatal: [localhost]: FAILED! => {"msg": "The task includes an option with an undefined variable. The error was: 'app_version' is undefined. 'app_version' is undefined\n\nThe error appears to be in '/runner/project/hello-world-playbook.yml': line 9, column 7"}
- Retrieved job events confirming failure

Active
- Presenting root cause and resolution steps to user (not yet done)

Blocked

```

- Now its upto the human to decide at what level of playbook/ inventory/host\_vars/group\_vars or extra variable the resolution should be proposed.

```

- Fix the failing AAP Job Template MCP-Error-Detection-Lab (ID 8) by adding app_version: "1.0" as an extra variable,

```

The Automation Orchestrator workflow can be integrated to provide a human-in-the-loop validation and proposed solution by the AI where with simple click of a button the Human can approve the fix and the system will automatically apply the fix and retry the playbook.

- I will add the Variable to the Job Template extra variable section as shown below and click on the Save button and execute the Job Template again. The playbook should now succeed without any errors:

Templates > Edit MCP-Error-Detection-Lab

### Edit MCP-Error-Detection-Lab

**Job type** ? ☐ Prompt on launch **Inventory** ? ☐ Prompt on launch **Project** ?

Run Demo Inventory MCP-Playbooks-Lab

**Playbook** ? ☐ Prompt on launch **Execution environment** ? ☐ Prompt on launch **Credentials** ?

hello-world-playbook.yml Select execution environment Select credentials

**Policy enforcement** ? ☐ Prompt on launch **Labels** ? ☐ Prompt on launch **Forks** ?

Enter policy enforcement Select or create labels 0

**Limit** ? ☐ Prompt on launch **Verbosity** ? ☐ Prompt on launch **Job slicing** ?

Enter limit to reduce number of hosts 0 (Normal) 1

**Timeout** ? ☐ Prompt on launch **Show changes** ? ☐ Prompt on launch **Instance groups** ?

0 Select instance groups

**Job tags** ? ☐ Prompt on launch **Skip tags** ? ☐ Prompt on launch

Select or create job tags Select or create skip tags

**Extra variables** ? ☒ Prompt on launch

app\_version: success

**Options**

☐ Privilege escalation ☐ Provisioning callback ☐ Enable webhook ☐ Concurrent jobs ☐ Enable fact storage ☐ Prevent instance group fallback

Save job template Cancel

## Watch Playbook Retry and Succeed

### Playbook Retries with Fix

The playbook automatically retries:

```
PLAY [Hello World Sample] ****

TASK [Hello Message] ****
ok: [localhost] => {
  "msg": "Hello World"
}

TASK [APP version] ****
ok: [localhost] => {
  "msg": "success"
}

PLAY RECAP ****
localhost: ok=2 changed=0 unreachable=0 failed=0

✓ PLAYBOOK COMPLETED SUCCESSFULLY!
```

+ image::1.10.png[]

+ image::1.11.png[]

### Alternate fix option:

The applied fix (visible in job output or GitHub):

```
# Original (with error):
- name: APP version
  debug:
    msg: "{{ app_version }}"

# Fixed by Claude (auto-applied):
- name: APP version
  debug:
    msg: "{{ app_version | default('1.0.0') }}"

# Or with explicit variable:
vars:
  app_version: "1.0.0"
```

▣ **Playbook now has the fix applied!**

# Hands-On Exercises & Summary

Try these exercises to practice:

## Exercise 1: Simple Query

Goal: Get a count of total jobs  
Question: "How many jobs exist in my AAP?"  
Expected: A number with job details  
Learns: Basic job querying

## Exercise 2: Filtered Query

Goal: Find specific jobs by status  
Question: "Show me only successful jobs from the past week"  
Expected: Filtered list of successful jobs  
Learns: Query with conditions

## Exercise 3: Multi-Step Query

Goal: Get multiple types of information  
Question: "How many jobs failed this week and what were their error messages?"  
Expected: List of failed jobs with error details  
Learns: MCP can do multi-step queries

## Exercise 4: Inventory Query

Goal: Understand your infrastructure  
Question: "Show me all hosts grouped by operating system"  
Expected: Breakdown of hosts by OS (Windows, Linux, etc)  
Learns: Inventory management queries

## Exercise 5: Follow-Up Conversation

Goal: Have a multi-turn conversation  
Exchange:  
You: "List my execution nodes"  
AI: [Returns list of nodes]  
You: "Which one has the most capacity?"  
AI: [Analyzes and compares]  
You: "What's the health status of that node?"  
AI: [Queries monitoring data]  
Learns: Contextual conversation

## Part 7: Advanced Tips

### Tip 1: Be Specific

#### Instead of:

"Show me jobs"

#### Say:

"Show me jobs that failed in the past 7 days with error details"

Specific queries get better answers.

### Tip 2: Ask Follow-Up Questions

You: "What's the status of my platform?"  
AI: [Gives health status]  
You: "Which execution node is having issues?"  
AI: [Digs deeper based on context]

The AI remembers context in the conversation.

### Tip 3: Combine Information

You: "I have failed jobs and offline hosts. Are they related?"  
AI: [Analyzes both, shows correlation]

The AI can synthesize information from multiple queries.

### Tip 4: Use Natural Language

MCP isn't a CLI. You don't need exact syntax. Just ask naturally:

"How's my system doing?"  
"Are my servers healthy?"  
"What's been happening lately?"  
"Do I have any problems?"

The AI will interpret and call the right tools.

## Part 8: When to Use MCP vs Other Methods

### Use MCP for:

- **Quick information** - "How many hosts do I have?"
- **Status checks** - "Is my platform healthy?"
- **Trend analysis** - "Have failures increased?"
- **Natural language** - Questions, not commands
- **Exploration** - "Show me everything about..."
- **Error Detection** - "Tell me why the sample-JT-name failed"

### Use AAP Web UI for:

- **Detailed editing** - Configuring large projects
- **Visual dashboards** - Seeing graphs and charts
- **Bulk operations** - Multi-host changes
- **Training new users** - Learning the UI

### Use AAP API for:

- **Automation** - Scripts and playbooks
- **Integration** - Custom tools
- **Exact control** - Precise parameter passing
- **Batch operations** - 1000s of changes

## Summary Checklist

By now, you should have:

- ☐ Run your first successful query against AAP
- ☐ Seen MCP tools being invoked
- ☐ Tried queries from multiple toolsets
- ☐ Had multi-turn conversations
- ☐ Encountered and solved at least one issue
- ☐ Understood what tools are available

## Congratulations!

You've completed all three labs:

1. ☐ Deployed MCP on OCP
2. ☐ Configured your AI tool

3. □ Ran queries and practiced troubleshooting.

**You now have a working MCP setup!**

# Appendix & Resources

This section provides quick references, checklists, and deep-dive resources for MCP and AAP integration.

## What's in the Appendix?

- **MCP Cheat Sheet** — Quick reference for common tasks and queries
- **Knowledge Check** — Test your understanding with a quiz
- **Troubleshooting Guide** — Solutions for complex issues
- **Useful Links** — Official documentation and references

## Quick Navigation

- [MCP Cheat Sheet](#) — Copy-paste queries and configurations
- [Knowledge Check & Quiz](#) — Verify you learned the concepts
- [Troubleshooting Guide](#) — Detailed debugging steps

## External Resources

### Official Documentation

- [Model Context Protocol Specification](#) — The MCP standard
- [Red Hat AAP MCP Server Documentation](#) — Official AAP MCP guide
- [Ansible Documentation](#) — General Ansible reference

### AI Tools

- [Cursor IDE MCP Documentation](#)
- [Claude Code Documentation](#)
- [LibreChat Documentation](#)

### Community & Support

- [Ansible Community Forum](#) — Ask questions and get help
- [Ansible TMM MCP Demo GitHub](#) — Real-world examples
- [MCP GitHub Organization](#) — MCP specs and tools

## How to Use This Appendix

1. **During the course:** Use the cheat sheet when you're in the labs

2. **After the course:** Reference troubleshooting when issues arise
3. **Learning verification:** Take the knowledge check to test yourself
4. **Advanced topics:** Links to deep-dive resources for specific areas

Good luck with your MCP journey!

## MCP Cheat Sheet

A quick reference for common tasks, queries, and configurations.

### Quick Setup Commands

#### Cursor IDE (One-Time Setup)

```
# 1. Add to ~/.zshrc or ~/.bashrc
export AAP_MCP_URL="https://your-aap.com:8448"
export MY_SERVICE_TOKEN="your-token-here"
export NODE_TLS_REJECT_UNAUTHORIZED=0

# 2. Reload shell
source ~/.zshrc

# 3. Open Cursor Settings → Tools & MCP
# 4. Add 6 servers with this template:
{
  "type": "streamable-http",
  "url": "${env:AAP_MCP_URL}/job_management/mcp",
  "headers": {"Authorization": "Bearer ${env:MY_SERVICE_TOKEN}"}
}
# (Change "job_management" to: inventory_management, system_monitoring,
#  user_management, security_compliance, platform_configuration)

# 5. Restart Cursor
```

#### Claude Desktop (One-Time Setup)

```
# 1. Add to ~/.zshrc or ~/.bashrc
export AAP_SERVER="https://your-aap.com:8448"
export AAP_SERVICE_TOKEN="your-token-here"
export NODE_TLS_REJECT_UNAUTHORIZED=0

# 2. Reload shell
source ~/.zshrc

# 3. Create/edit ~/.claude/claude.json with 6 servers
# (See Lab 2 for full config)
```

## # 4. Restart Claude

### Test Configuration

```
# Verify MCP server is accessible
curl -k -H "Authorization: Bearer $MY_SERVICE_TOKEN" \
  https://your-aap.com:8448/job_management/mcp

# Verify token is set
echo $MY_SERVICE_TOKEN
echo $AAP_MCP_URL

# Verify environment is loaded
env | grep AAP
```

### Common Queries

#### Job Management Queries

| What You Want      | What to Ask                                     |
|--------------------|-------------------------------------------------|
| List all jobs      | "List all job executions"                       |
| Recent failed jobs | "Show me jobs that failed in the last 7 days"   |
| Running jobs       | "What jobs are currently running?"              |
| Job details        | "Show me details for job #456"                  |
| Job by name        | "Did the 'Deploy Web Servers' job succeed?"     |
| Job count          | "How many jobs have run this month?"            |
| Job templates      | "List all job templates available"              |
| Job output         | "Show me the output from the last run of Job X" |

#### Inventory Queries

| What You Want      | What to Ask                                    |
|--------------------|------------------------------------------------|
| Host count         | "How many hosts are in my inventory?"          |
| Hosts by group     | "Which hosts are in the production group?"     |
| Windows hosts      | "Show me all Windows servers"                  |
| Linux hosts        | "List all Linux hosts"                         |
| Offline hosts      | "Are there any hosts that haven't checked in?" |
| Host variables     | "Show me variables for host web-01"            |
| Groups list        | "What inventory groups do I have?"             |
| Specific inventory | "Show me all hosts in the 'prod' inventory"    |

## Monitoring Queries

| What You Want   | What to Ask                                      |
|-----------------|--------------------------------------------------|
| Platform health | "What's the health status of my AAP instance?"   |
| Execution nodes | "Show me my execution nodes and their status"    |
| Capacity        | "What's my AAP capacity utilization?"            |
| Job metrics     | "How many jobs ran this week?"                   |
| Failure rate    | "What's my job failure percentage?"              |
| Performance     | "Which execution node has the best performance?" |

## User Management Queries

| What You Want | What to Ask                           |
|---------------|---------------------------------------|
| All users     | "List all users in the system"        |
| Teams         | "What teams exist in AAP?"            |
| Team members  | "Show me the members of the ops team" |
| User roles    | "What roles does user alice have?"    |
| Organizations | "List all organizations"              |
| User count    | "How many users are in my AAP?"       |

## Security & Compliance Queries

| What You Want   | What to Ask                                      |
|-----------------|--------------------------------------------------|
| Audit logs      | "Show me audit logs from the past 24 hours"      |
| Credentials     | "List all credentials and their types"           |
| Changes         | "What changed in the past week?"                 |
| Access events   | "Show me login activity"                         |
| Failed attempts | "Were there any failed authentication attempts?" |

## Configuration Queries

| What You Want          | What to Ask                                  |
|------------------------|----------------------------------------------|
| Projects               | "List all projects and their status"         |
| Execution environments | "What execution environments are available?" |
| Settings               | "Show me platform configuration"             |
| Integrations           | "What integrations are configured?"          |
| Database               | "Show me database settings"                  |

# API Token Quick Reference

## Create Token (AAP Web UI)

1. Log in to AAP
  2. Access Management → API Tokens
  3. Create API Token
  4. Name: "mcp-token" (or any name)
  5. Application: Leave blank
  6. Scope: "Read" or "Read + Write"
  7. Click Create
  8. Copy token immediately (not shown again!)

## Token Scopes

| Scope        | Permissions                      |
|--------------|----------------------------------|
| Read         | View-only access (safe default)  |
| Write        | Can make changes (use carefully) |
| Read + Write | Full access (admin only)         |

## Rotate Token (Security Best Practice)

1. Create new token with same scope
  2. Update configuration in AI tool
  3. Restart AI tool
  4. Delete old token in AAP

## URL Reference

### Deployment Types

| Deployment           | URL Format                                                                                      |
|----------------------|-------------------------------------------------------------------------------------------------|
| OpenShift            | <a href="https://aap-mcp.apps.cluster.example.com">https://aap-mcp.apps.cluster.example.com</a> |
| Containerized (RHEL) | <a href="https://aap-hostname.com:8448">https://aap-hostname.com:8448</a>                       |

### MCP Endpoint Paths

| Toolset              | Endpoint                                                          |
|----------------------|-------------------------------------------------------------------|
| Job Management       | <a href="/job_management/mcp">/job_management/mcp</a>             |
| Inventory Management | <a href="/inventory_management/mcp">/inventory_management/mcp</a> |
| System Monitoring    | <a href="/system_monitoring/mcp">/system_monitoring/mcp</a>       |
| User Management      | <a href="/user_management/mcp">/user_management/mcp</a>           |

| Toolset                | Endpoint                    |
|------------------------|-----------------------------|
| Security & Compliance  | /security_compliance/mcp    |
| Platform Configuration | /platform_configuration/mcp |

## Full MCP URLs

```
# Example: OpenShift deployment
https://aap-mcp.apps.example.com/job_management/mcp
https://aap-mcp.apps.example.com/inventory_management/mcp
https://aap-mcp.apps.example.com/system_monitoring/mcp
https://aap-mcp.apps.example.com/user_management/mcp
https://aap-mcp.apps.example.com/security_compliance/mcp
https://aap-mcp.apps.example.com/platform_configuration/mcp

# Example: Containerized deployment
https://aap.example.com:8448/job_management/mcp
https://aap.example.com:8448/inventory_management/mcp
(etc...)
```

## Troubleshooting Checklist

### Connection Issues

| Check               | How to Verify                                   |
|---------------------|-------------------------------------------------|
| URL is correct      | Paste URL in browser, should not give "refused" |
| Token is valid      | Try creating new token in AAP                   |
| Network reaches MCP | <code>curl -k https://your-mcp-url</code>       |
| MCP is deployed     | Check OCP: Deployments → mcp pod running?       |
| AAP is accessible   | Log in to AAP web UI successfully               |

### Configuration Issues

| Check                     | How to Verify                                                                |
|---------------------------|------------------------------------------------------------------------------|
| Environment variables set | <code>echo \$MY_SERVICE_TOKEN</code>                                         |
| All 6 toolsets added      | "What tools are available?" should list 85+                                  |
| Bearer prefix included    | Config should have <code>Bearer \$token</code> not just <code>\$token</code> |
| AI tool restarted         | Fully close and reopen (not refresh)                                         |
| YAML is valid             | Try different tool first to test setup                                       |

### Permission Issues

| Check                     | How to Verify                            |
|---------------------------|------------------------------------------|
| Token scope matches needs | Read-only token can't launch jobs        |
| AAP user has permissions  | Check RBAC in AAP → Access Management    |
| Organization access       | User should be in organization to access |
| Team membership           | Team-level resources require team access |

## Quick Debugging Steps

### "Connection Refused"

```
# Step 1: Test URL manually
curl -k https://your-mcp-url
# Should return 400/401, not "refused"

# Step 2: Check deployment
oc get deployments -n aap | grep mcp
# Should show 1/1 Ready

# Step 3: Check route
oc get routes -n aap | grep mcp
# Should show a URL in Location column

# Step 4: Try different URL format
# If using OCP, make sure NO port (aap-mcp.apps.com)
# If using RHEL, make sure WITH port (hostname:8448)
```

### "Unauthorized" / "401"

```
# Step 1: Verify token
echo $MY_SERVICE_TOKEN
# Should show 40+ character string

# Step 2: Test token
curl -k -H "Authorization: Bearer $MY_SERVICE_TOKEN" \
  https://your-mcp-url
# Should not return 401

# Step 3: Create new token
# Go to AAP → Access Management → API Tokens
# Create new token
# Update configuration

# Step 4: Restart tool
# Fully close and reopen AI tool
```

## "Tool Not Found"

```
# Step 1: List available tools
# In AI tool, ask: "What tools are available?"

# Step 2: Verify all 6 toolsets configured
# Check config for all 6 endpoints

# Step 3: Check tool name spelling
# Use exact names from "tools available" response

# Step 4: Restart tool and retry
```

# Common Configuration Errors

## Missing Bearer Prefix

### Wrong:

```
"Authorization": "${env:MY_SERVICE_TOKEN}"
```

### Correct:

```
"Authorization": "Bearer ${env:MY_SERVICE_TOKEN}"
```

## Wrong URL Format

### Wrong (OpenShift):

```
https://aap-mcp.apps.example.com:8448/job_management/mcp
(includes port)
```

### Correct (OpenShift):

```
https://aap-mcp.apps.example.com/job_management/mcp
(no port)
```

## Missing Toolsets

### Wrong (only 1 server):

```
{
  "aap-job-mgmt": {...}
}
```

```
}
```

**Correct (all 6 servers):**

```
{
  "aap-job-mgmt": {...},
  "aap-inventory": {...},
  "aap-monitoring": {...},
  "aap-user-mgmt": {...},
  "aap-security": {...},
  "aap-config": {...}
}
```

## Environment Variables

| Variable                              | Purpose                            | Example                                                                         |
|---------------------------------------|------------------------------------|---------------------------------------------------------------------------------|
| AAP_MCP_URL or AAP_SERVER             | MCP server location                | <a href="https://aap-mcp.apps.example.com">https://aap-mcp.apps.example.com</a> |
| MY_SERVICE_TOKEN or AAP_SERVICE_TOKEN | API authentication                 | abc123def456...                                                                 |
| NODE_TLS_REJECT_UNAUTHORIZED          | Skip SSL verification (demo only!) | 0                                                                               |
| NODE_DEBUG                            | Debug mode (verbose logging)       | *                                                                               |

**Warning:** Only use `NODE_TLS_REJECT_UNAUTHORIZED=0` in development! Use proper SSL certificates in production.

## Useful One-Liners

```
# Find MCP route URL
oc get routes -n aap -o wide | grep mcp

# Get MCP pod logs
oc logs -n aap -l app=aap-mcp --tail=50

# Check MCP pod status
oc get pods -n aap | grep mcp

# Restart MCP pod
oc delete pod -n aap -l app=aap-mcp

# Check AAP token validity
curl -k -H "Authorization: Token $MY_SERVICE_TOKEN" \
  https://your-aap-url/api/v2/ping/
```

```
# Test MCP endpoint
curl -k -H "Authorization: Bearer $MY_SERVICE_TOKEN" \
  https://your-mcp-url/job_management/mcp

# Show environment variables
env | grep -E "(AAP|TOKEN|MCP)"
```

## Quick Reference Table

| Component               | Where         | File/Location                  |
|-------------------------|---------------|--------------------------------|
| Course Content          | This site     | Start at Welcome               |
| MCP Server              | OpenShift     | Deployments → aap-mcp pod      |
| MCP Server URL          | OpenShift     | Networking → Routes → aap-mcp  |
| API Token               | AAP Web UI    | Access Management → API Tokens |
| Tool Config (Cursor)    | Computer      | Cursor Settings → Tools & MCP  |
| Tool Config (Claude)    | Computer      | ~/.claude/claude.json          |
| Tool Config (LibreChat) | Web Interface | Settings → MCP Servers         |
| Environment Variables   | Computer      | ~/.zshrc or ~/.bashrc          |
| AAP User Permissions    | AAP Web UI    | Access Management → Users      |

**Pro Tip:** Bookmark this page for quick reference during labs and production use!

## Knowledge Check & Quiz

Test your understanding of MCP and its integration with Ansible Automation Platform.

This is a self-assessment quiz. Answer the questions honestly to see what you've learned and what areas need review.

### Part 1: MCP Fundamentals (5 Questions)

#### Question 1: What is MCP?

Which best describes the Model Context Protocol?

a) A method for encrypting API tokens b) A standard for connecting AI models to external systems

✓ c) A replacement for Ansible playbooks d) A database technology

**Answer:** B

**Explanation:** MCP (Model Context Protocol) is an open-source standard that provides a universal way for AI models to interact with external data and tools. It's like "USB-C for the AI ecosystem."

### Question 2: Why Does MCP Matter for AAP?

What problem does MCP solve for Ansible users?

a) Makes playbooks run faster b) Reduces the cost of AAP licenses c) Gives AI real-time access to live AAP data ✓ d) Eliminates the need for SSH

**Answer:** C

**Explanation:** AI models are trained on frozen historical data. MCP solves this by giving them access to **real-time** data from your AAP environment—what jobs ran today, current host status, live inventory data, etc.

### Question 3: How Many Components Are in MCP Architecture?

Which of these is NOT a core MCP component?

a) AI Client (where you interact) b) AI Model (the LLM) c) MCP Server (the gatekeeper) d) Load Balancer ✓

**Answer:** D

**Explanation:** The four core components are: AI Client, AI Model, MCP Server, and Your System (AAP). A load balancer is infrastructure, not part of the core MCP architecture.

### Question 4: What's the MCP Server's Main Job?

What does the MCP Server do?

a) Trains AI models b) Validates requests and enforces security ✓ c) Stores your AAP data d) Runs your Ansible playbooks

**Answer:** B

**Explanation:** The MCP Server is the "gatekeeper." It authenticates requests using your API token, validates that the request is allowed based on RBAC, translates the request to an AAP API call, and standardizes the response.

### Question 5: Why is Authentication Important in MCP?

What would happen if MCP didn't enforce authentication?

a) It would be faster but less secure ✓ b) Nothing different c) Playbooks would execute automatically d) It would require you to log in twice

**Answer:** A

**Explanation:** Without authentication, anyone could use your MCP server, which would be a severe security risk. Authentication ensures only authorized users can access your infrastructure through MCP.

---

## Part 2: AAP Toolsets (6 Questions)

### Question 6: How Many AAP MCP Toolsets Are There?

How many separate toolsets does the Ansible MCP Server expose?

a) 3 b) 4 c) 6 ✓ d) 12

**Answer:** C

**Explanation:** The 6 AAP MCP toolsets are: Job Management, Inventory Management, System Monitoring, User Management, Security & Compliance, and Platform Configuration.

### Question 7: Which Toolset Would You Use to Find Offline Hosts?

If you want to find which hosts haven't checked in recently, which toolset would you use?

a) Job Management b) Inventory Management ✓ c) System Monitoring d) Security & Compliance

**Answer:** B

**Explanation:** The Inventory Management toolset handles hosts, groups, and inventory queries. It would allow you to query which hosts are in your inventory and their status.

### Question 8: What Does the Security & Compliance Toolset Provide?

Which of these is NOT a capability of the Security & Compliance toolset?

a) View audit logs b) List credentials c) Launch job templates ✓ d) Review access history

**Answer:** C

**Explanation:** Launching job templates is handled by the Job Management toolset. Security & Compliance provides access to audit logs, credentials, security policies, and access history.

### Question 9: Can You Use All 6 Toolsets with One MCP Server?

Can you connect to multiple toolsets simultaneously?

a) No, you can only use one at a time b) Yes, but only if you have admin permissions c) Yes, all 6 can be configured together ✓ d) No, you need separate MCP servers for each

**Answer:** C

**Explanation:** You configure all 6 toolsets when setting up your AI tool. The single MCP server has 6 endpoints, and you connect to all of them. The AI model can call any tool needed.

**Question 10: Which Toolset Gives System Health Status?**

If you want to check "Is my AAP platform healthy?", which toolset would you ask?

a) User Management b) System Monitoring ✓ c) Platform Configuration d) Security & Compliance

**Answer:** B

**Explanation:** System Monitoring provides health status, capacity metrics, and performance data about your AAP instance and execution nodes.

**Question 11: What's the Difference Between Job Management and Inventory Management?**

What does Job Management handle that Inventory Management doesn't?

a) Both handle hosts b) Job Management is for playbooks, Inventory is for servers c) Job Management handles job execution, workflows, and templates; Inventory handles hosts and groups ✓ d) They're the same, just different names

**Answer:** C

**Explanation:** Job Management deals with automation workflows—running jobs, checking status, accessing job templates. Inventory Management deals with infrastructure—hosts, groups, variables.

---

## Part 3: Practical Workflow (5 Questions)

**Question 12: What Three Steps Deploy and Use MCP?**

What are the three main steps to get MCP working?

a) Download, Install, Verify b) Deploy MCP server, Generate token, Configure tool ✓ c) Login, Setup, Run d) Cluster, Pod, Service

**Answer:** B

**Explanation:** The workflow is: (1) Deploy the MCP server on your infrastructure, (2) Generate an API token in AAP, (3) Configure your AI tool with the token and MCP URL.

**Question 13: Where Does the API Token Come From?**

Where do you create an API token?

a) In Cursor IDE settings b) In the Claude configuration file c) In Ansible Automation Platform (AAP) ✓ d) In OpenShift console

**Answer:** C

**Explanation:** API tokens are created in AAP under Access Management → API Tokens. You then use that token to configure your AI tool.

#### Question 14: What Information Do You Need to Configure an AI Tool?

When configuring Cursor or Claude for MCP, what two pieces of information are essential?

- a) Username and password b) AAP hostname and cluster name c) MCP Server URL and API Token ✓
- d) Playbook path and inventory file

**Answer:** C

**Explanation:** You need to tell the AI tool: (1) Where is the MCP server? (the URL), (2) Who am I? (the API token). The tool uses these to authenticate and connect.

#### Question 15: What Happens if Your API Token Expires?

If your API token expires, what do you need to do?

- a) Reinstall AAP b) Recreate the token and update your tool configuration ✓ c) Wait for it to auto-renew d) Upgrade Cursor or Claude

**Answer:** B

**Explanation:** Tokens can expire or be revoked. When this happens, you create a new token in AAP and update your AI tool configuration with the new token, then restart the tool.

#### Question 16: What's the Difference Between Read and Write Token Scope?

What's the key difference in what you can do with each scope?

- a) Read is slower, Write is faster b) Read is for Cursor, Write is for Claude c) Read allows querying but not making changes; Write allows both ✓ d) There's no difference

**Answer:** C

**Explanation:** A "Read" scoped token can only query information (view jobs, list hosts, check status). A "Write" scoped token can also make changes (launch jobs, create users, modify settings). Start with Read for safety.

---

## Part 4: Security & Architecture (4 Questions)

#### Question 17: Where Does Security Get Enforced?

At what point does MCP enforce your AAP permissions?

- a) Only in your AI tool b) Only in the MCP Server ✓ c) Only in AAP d) The AI model enforces it

**Answer:** B

**Explanation:** The MCP Server is the security checkpoint. It validates your token, checks your RBAC permissions, and only allows actions your user account has access to. The AI model doesn't have direct access to AAP.

### Question 18: Can the AI Model Bypass Your Permissions?

If your token is read-only, can an AI tool still launch jobs?

a) Yes, if the AI is smart enough b) Yes, but only with admin help c) No, the MCP Server will reject the request ✓ d) Only if you ask it politely

**Answer:** C

**Explanation:** The MCP Server enforces permissions at the security checkpoint. No matter how the AI requests it, if the token doesn't have write access, the request is denied.

### Question 19: What Gets Sent Between Components?

What information flows from the AI to the MCP Server?

a) Your AAP password b) Structured tool calls with parameters ✓ c) Ansible playbooks d) SSH private keys

**Answer:** B

**Explanation:** The AI sends structured tool calls ("call job\_list with limit\_days=7"), not raw requests or credentials. The MCP Server then translates these to API calls with your token.

### Question 20: If You Delete Your API Token, What Happens?

What happens to existing MCP connections if you delete the token?

a) They keep working (tokens can't be deleted) b) They fail immediately ✓ c) They keep working for 24 hours then fail d) Nothing, tokens don't affect connections

**Answer:** B

**Explanation:** When you delete a token in AAP, all connections using that token are immediately invalidated. You'll get "Unauthorized" errors until you create a new token and reconfigure your tool.

---

## Part 5: Real-World Scenarios (5 Questions)

### Question 21: Scenario - Your First Query

You've configured Cursor with MCP. You ask: "Show me failed jobs from this week."

What happens?

a) Cursor searches the internet b) Cursor asks the AI model, which calls MCP tools, which queries AAP ✓ c) Cursor runs a local command d) Nothing—you need to manually query AAP first

**Answer:** B

**Explanation:** The full workflow: You ask → Cursor sends to AI → AI interprets and calls MCP tools → MCP Server queries AAP → Data comes back → AI summarizes for you.

### Question 22: Scenario - Permission Error

You have a read-only token and try to ask: "Launch the deploy job."

What happens?

a) It launches the job b) The MCP Server returns a permission denied error ✓ c) The job launches but is marked as "read-only" d) It asks you for a password

**Answer:** B

**Explanation:** The MCP Server checks your token scope. Read-only tokens cannot launch jobs. The request is rejected at the security checkpoint.

### Question 23: Scenario - Connection Troubleshooting

Cursor can't connect to MCP. You've verified the URL and token are correct. What's a likely cause?

a) You need to upgrade Cursor b) The MCP deployment pod isn't running ✓ c) You configured too many toolsets d) The API token is the wrong color

**Answer:** B

**Explanation:** If the URL and token are correct but connection is refused, the MCP server itself is probably not running. Check that the MCP deployment is actually deployed and the pod is ready.

### Question 24: Scenario - Best Practice

For production use with a team, what's the recommended approach for API tokens?

a) Share one token among all team members b) Create individual tokens for each team member ✓ c) Use the admin token for everything d) Rotate tokens weekly

**Answer:** B

**Explanation:** Create individual tokens so you can track who did what (audit logs), revoke access independently, and follow the principle of least privilege. One shared token is a security risk.

### Question 25: Scenario - Adding Another Tool

You want to use both Cursor and Claude with the same AAP instance. What do you need?

a) Two different AAP instances b) Two different MCP servers c) One token for each tool (or can share same token) ✓ d) Different AAP users for each tool

**Answer:** C

**Explanation:** You can use the same API token for multiple AI tools, or create separate tokens for each. Either works. Both tools connect to the same MCP server(s).

## Answer Key & Scoring

### Scoring

| Score    | Assessment                                              |
|----------|---------------------------------------------------------|
| 25/25    | Expert! You've mastered MCP.                            |
| 22-24/25 | Advanced. You understand the concepts deeply.           |
| 18-21/25 | Proficient. You're ready for production use.            |
| 15-17/25 | Intermediate. Review sections where you missed.         |
| 12-14/25 | Beginner. Take the course again focusing on weak areas. |
| <12/25   | Review fundamentals. Go back to Chapter 1.              |

### By Topic

**MCP Fundamentals (Q1-5):** - If you scored 5/5: Great understanding of what MCP is - If you scored <4/5: Review Section 1-2 of Chapter 1

**AAP Toolsets (Q6-11):** - If you scored 6/6: Perfect knowledge of toolsets - If you scored <5/6: Review Section 3 of Chapter 1

**Practical Workflow (Q12-16):** - If you scored 5/5: Ready for Labs - If you scored <4/5: Review Section 5 of Chapter 1

**Security & Architecture (Q17-20):** - If you scored 4/4: Security understanding is solid - If you scored <3/4: Review Section 2 and 4 of Chapter 1

**Real-World Scenarios (Q21-25):** - If you scored 5/5: Apply knowledge confidently - If you scored <4/5: Practice more before production

## What to Review Based on Results

### If You Struggled with Fundamentals (Q1-5)

You might not fully understand what MCP is or why it matters. Review: \* Section 1: What is MCP? \* Section 2: MCP Architecture

Focus on understanding: \* MCP is a **standard**, not a tool \* It gives AI **real-time access** to your systems \* The MCP Server is the **security checkpoint**

### If You Struggled with Toolsets (Q6-11)

You might not know which toolset to use for different tasks. Review: \* Section 3: AAP Integration & Toolsets

Practice: \* Map real tasks to toolsets (e.g., "find offline hosts" → Inventory Management) \* Ask yourself "which toolset would handle this?" for different scenarios

### If You Struggled with Practical Workflow (Q12-16)

You might not understand the deployment and configuration steps. Review: \* Section 5: End-to-End Workflow \* Lab 1, 2, 3 for hands-on practice

Practice: \* Run through the labs again \* Focus on each step and understand **why** you're doing it

### If You Struggled with Security (Q17-20)

You might not understand how permissions and authentication work. Review: \* Section 2: MCP Architecture (component roles) \* Section 4: Communication Workflow (security checkpoints)

Key concepts: \* AI can't access AAP directly; MCP Server is the gatekeeper \* Tokens control what you can do \* RBAC is enforced at the server, not the client

### If You Struggled with Real-World Scenarios (Q21-25)

You might have trouble applying concepts to realistic situations. Review: \* All previous sections (understand foundations first) \* Lab 3: Run Queries and Troubleshoot

Practice: \* Set up MCP yourself and run queries \* Troubleshoot real issues \* Think through "what would happen if..." scenarios

---

## Tips for Using This Knowledge Check

1. **Take it honestly** — This is for you, not graded
2. **Review weaknesses** — Use the scoring guide to focus study
3. **Re-take it later** — Test yourself again after time passes
4. **Apply it** — The real test is using MCP successfully
5. **Help others** — Teaching others deepens your understanding

---

## Passing the Course

To successfully complete this MCP course, you should be able to:

□ Define MCP and explain why it matters

☐ Describe the 4 MCP components and how they interact

☐ Name the 6 AAP toolsets and when to use each

☐ Deploy MCP on OpenShift

☐ Generate an API token securely

☐ Configure a tool (Cursor, Claude, or LibreChat)

☐ Run queries and understand the workflow

☐ Troubleshoot common issues

☐ Score 18+ on this knowledge check

If you meet these criteria, **congratulations!** You're ready to use MCP in your environment.

---

### Next Steps:

- Use the MCP Cheat Sheet for quick reference
- Reference the Troubleshooting Guide when issues arise
- Share what you've learned with your team
- Set up MCP in your production environment with confidence